

Tensorflow

Overview

TensorFlow uses a computational graph to represent a series of mathematical operations (nodes) and the flow of data between them (edges).

```
import tensorflow as tf

# Create tensors
tensor_1 = tf.constant([[1, 2, 3], [4, 5, 6]])
tensor_2 = tf.constant([[7, 8, 9], [10, 11, 12]])

# Create variables
variable_1 = tf.Variable(tensor_1)
variable_2 = tf.Variable(tensor_2)

# Perform operations on tensors and variables
sum_tensor = tf.add(tensor_1, tensor_2)
product_tensor = tf.multiply(tensor_1, tensor_2)
sum_variable = tf.add(variable_1, variable_2)
```

Advantages

- OpenSource and Active Community
- Flexible and scalable
- Extensive documentation and Resources
- Tensorboard for visualization
- TFX for production

Disadvantages

- No windows support
- Comparatively slow
- Only NVIDIA GPU support
- Architectural limitations
- Inconsistent

Extra terms

- Tensor:
 - The fundamental data type in Tensorflow. Represents multi-dimensional arrays
 - Can be constants or variables
 - Can have different ranks (dimensions)
 - Types
 - Scaler Tensor → rank-0
 - Vector Tensor → rank-1
 - Matrix Tensor → rank-2

- 3D and higher dimensional tensors → rank > 2
- Sparse Tensor → represents tensors with a large number of zero values
- Ragged tensor → tensors with non-uniform shapes, `tf.ragged.constant([..])`
- Variable Tensor → mutable tensors, `tf.Variable([1,2,3])`
- String Tensor → `tf.constant("hello world")`
- `tf.constant`: create a constant tensor
- `tf.Variable`: typically used for model parameters that need to be updated during the training
 - **Some methods**
- `tf.function` → decorator that converts a Python function to a TensorFlow graph function
- `tf.optimizers` → Contains various optimization functions for updating the model
- `tf.keras.models` → Module for defining and training ml models using high-level API
- `Tf.data` → Module for building scalable and efficient input pipelines for data preprocessing and augmentation
- `tf.keras.layers` → A module that provides a variety of pre-built layers for nn
- `tf.image` → Module for image processing operations such as resizing, cropping, etc
- `tf.reduce_mean`, `tf.reduce_sum` → Methods for reducing tensors along specified dimensions.
- `tf.losses` → Module that contains various loss functions
- `tf.saved_model.load()` → load the saved model in tf

Keras

Keras is a high-level neural networks API written in Python and integrated into TensorFlow.

```
import tensorflow as tf
from tensorflow.keras import Sequential
from tensorflow.keras.layers import Dense

# Create a simple feedforward neural network
model = Sequential([
    Dense(64, activation='relu', input_shape=(input_dimension,)),
    Dense(32, activation='relu'),
    Dense(output_dimension, activation='softmax')
])

# Display the model architecture
model.summary()
```

Why Keras?

- Simplicity and user-friendly interface
- Compatibility with TensorFlow
- Modularity and extensibility: can stack predefined building blocks together (layers)
- Gives consistency
- Integrated with TensorFlow tools such as TFX

Pytorch

Advantages

- *Dynamic Computational Graph: also known as eager execution, making easy to debug and experiment models*
- *Extensive community and research scope*
- *TorchScript for production*
- *Native support for Dynamic models → well suited for sequence-to-sequence models and RNN*
- *Libraries available for research → torch vision, torch audio*

Disadvantages

- *Deployment overhead, compared to tf*
- *Less Production-ready tools*
- *Smaller community compared to tf*
- *Not have enough deployment options for mobile and embedded devices*

Extra terms

- *Tensor → The fundamental data type in PyTorch, representing a multi-dimensional array*
 - `torch.Tensor([1,2,43])`

```
import torch

# Create tensors
tensor_1 = torch.tensor([[1, 2, 3], [4, 5, 6]])
tensor_2 = torch.tensor([[7, 8, 9], [10, 11, 12]])

# Create variables (requires_grad is set to True for tracking gradients)
variable_1 = torch.tensor(tensor_1, requires_grad=True)
variable_2 = torch.tensor(tensor_2, requires_grad=True)

# Perform operations on tensors and variables
sum_tensor = tensor_1 + tensor_2
product_tensor = torch.mul(tensor_1, tensor_2)
sum_variable = variable_1 + variable_2
```

Methods

- *torch.nn.Module → base class for all neural network model*

```
import torch.nn as nn


class MyModel(nn.Module):
    def __init__(self):
        super(MyModel, self).__init__()
        self.fc = nn.Linear(10, 5)
```

- *torch.nn.functional → a module containing a variety of functions that can be applied to tensors*
- *torch.nn.optim → module that contains various optimization algos*
- *torch.utils.data.Dataset → An abstract class representing a dataset in torch*


RNN

Is it a type of neural network architecture designed to process sequential data.

eg:-



So each of the RNN units has 3 weights and 2 biases. Bias will be there at the calculation of the output value and at the calculation of the hidden state.



Advantages

- Sequential information processing
- Flexibility in input and output length
- maintain a hidden state that captures information from previous inputs
- Natural representation of temporal patterns

Disadvantages

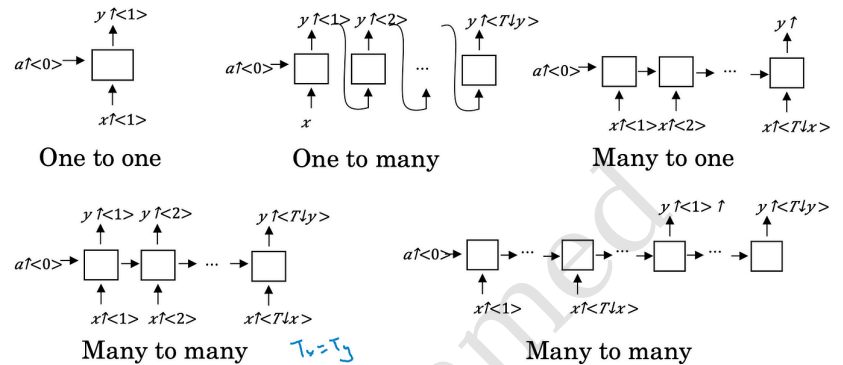
- Vanishing gradient → gradients become too small in training
- Exploding gradient → gradients become extremely high in training
- Difficulty in Learning Long-term dependencies
- Limited context
- Difficulty in irregular time intervals
- Difficulty in capturing global context

Extra terms

- Hidden state → a vector that represents the memory of the RNN, only in GRU and LSTM
- BPTT → Backpropagation through time

Types

- One-to-one → normal FNN
- One-to-Many → A single input time step is used to generate multiple output timesteps, for ex:-music generation
- Many-to-one → sentiment analysis
- Many-to-Many → sequence to sequence models



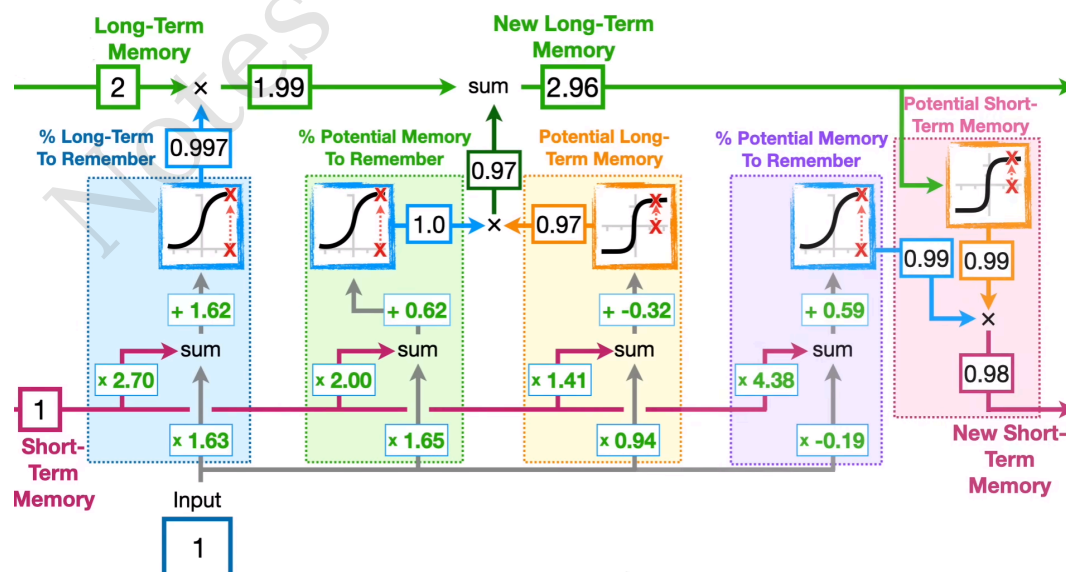
Andrew Ng

Evaluation metrics

- Accuracy, precision, recall, f1-score → classification
- MSE, RMSE, MAE, R2 → regression
- BLEU score (bilingual evaluation understudy) →
 - measures the quality of the machine-generated text compared to human-generated reference text.
 - Primarily focus on the n-gram precision
 - The BLEU score is between 0 and 1
 - BLEU = brevity penalty \times exp of cumulative log precision
- METEOR, ROUGE → similar to BLEU

LSTM

Designed to overcome the vanishing gradient problem in RNN



Terms

- Cell state → runs along the sequence, allowing information to be passed from one time step to the other
- Hidden state → responsible for maintaining and updating the context information
- Gates
 - Forget Gate → determines what information from the cell state should be discarded
 - Input Gate → determines what new information to store in the cell state
 - Output Gate → chooses what information to output based on the cell state

Advantages

- Handle long-term dependencies
- Gating mechanism
- Robust training
- Adaptability to various tasks

Disadvantages

- Complexity
- Overfitting → mainly in small datasets
- Training time will be higher
- Difficulty in interpretability

Why tanh ???

- Tanh's output range is (-1, 1)
- Don't have the vanishing gradient problem
- Underlying for the gating system in Lstm

Regularization

- **L2 regularization**
 - Also known as weight decay
 - It adds a penalty term to the loss function
 - Prevent overfitting

$$\text{loss} = \text{original loss} + \frac{\lambda}{2} \sum_i \|W_i\|_2^2$$

- **Batch normalization**
 - designed to address the internal covariate shift problem during training.
 - It normalizes the input of each layer by subtracting the mean and dividing it by the standard deviation.
 - Can increase stability

$$\text{Batch Normalization}(x) = \frac{x - \text{mean}(x)}{\sqrt{\text{variance}(x) + \epsilon}} \times \gamma + \beta$$

- **Dropout**
 - It randomly drops (sets to zero) a fraction of the input units during training.
 - It helps prevent overfitting by introducing redundancy in the network.

- Prevent overfitting

Loss fn

- Binary classify → Binary cross-entropy
- Multiclass → categorical cross-entropy
- Multiclass (when target values are integers rather than one-hot vectors) → sparse categorical cross-entropy
- Multilabel → Binary cross-entropy
- Regression → MSE, MAE, Huber loss

Optimizers

- Batch Gradient Descent → the traditional one
- SGD → divide into batches and do gradient descent
- RMSProp → Different learning rates for each
- Momentum → have a momentum oscillator
- Adam → momentum + RMSProp